Dual-GPU Evaluation Overview

Engine compatibility, context scaling, speculative decoding and deployment validation

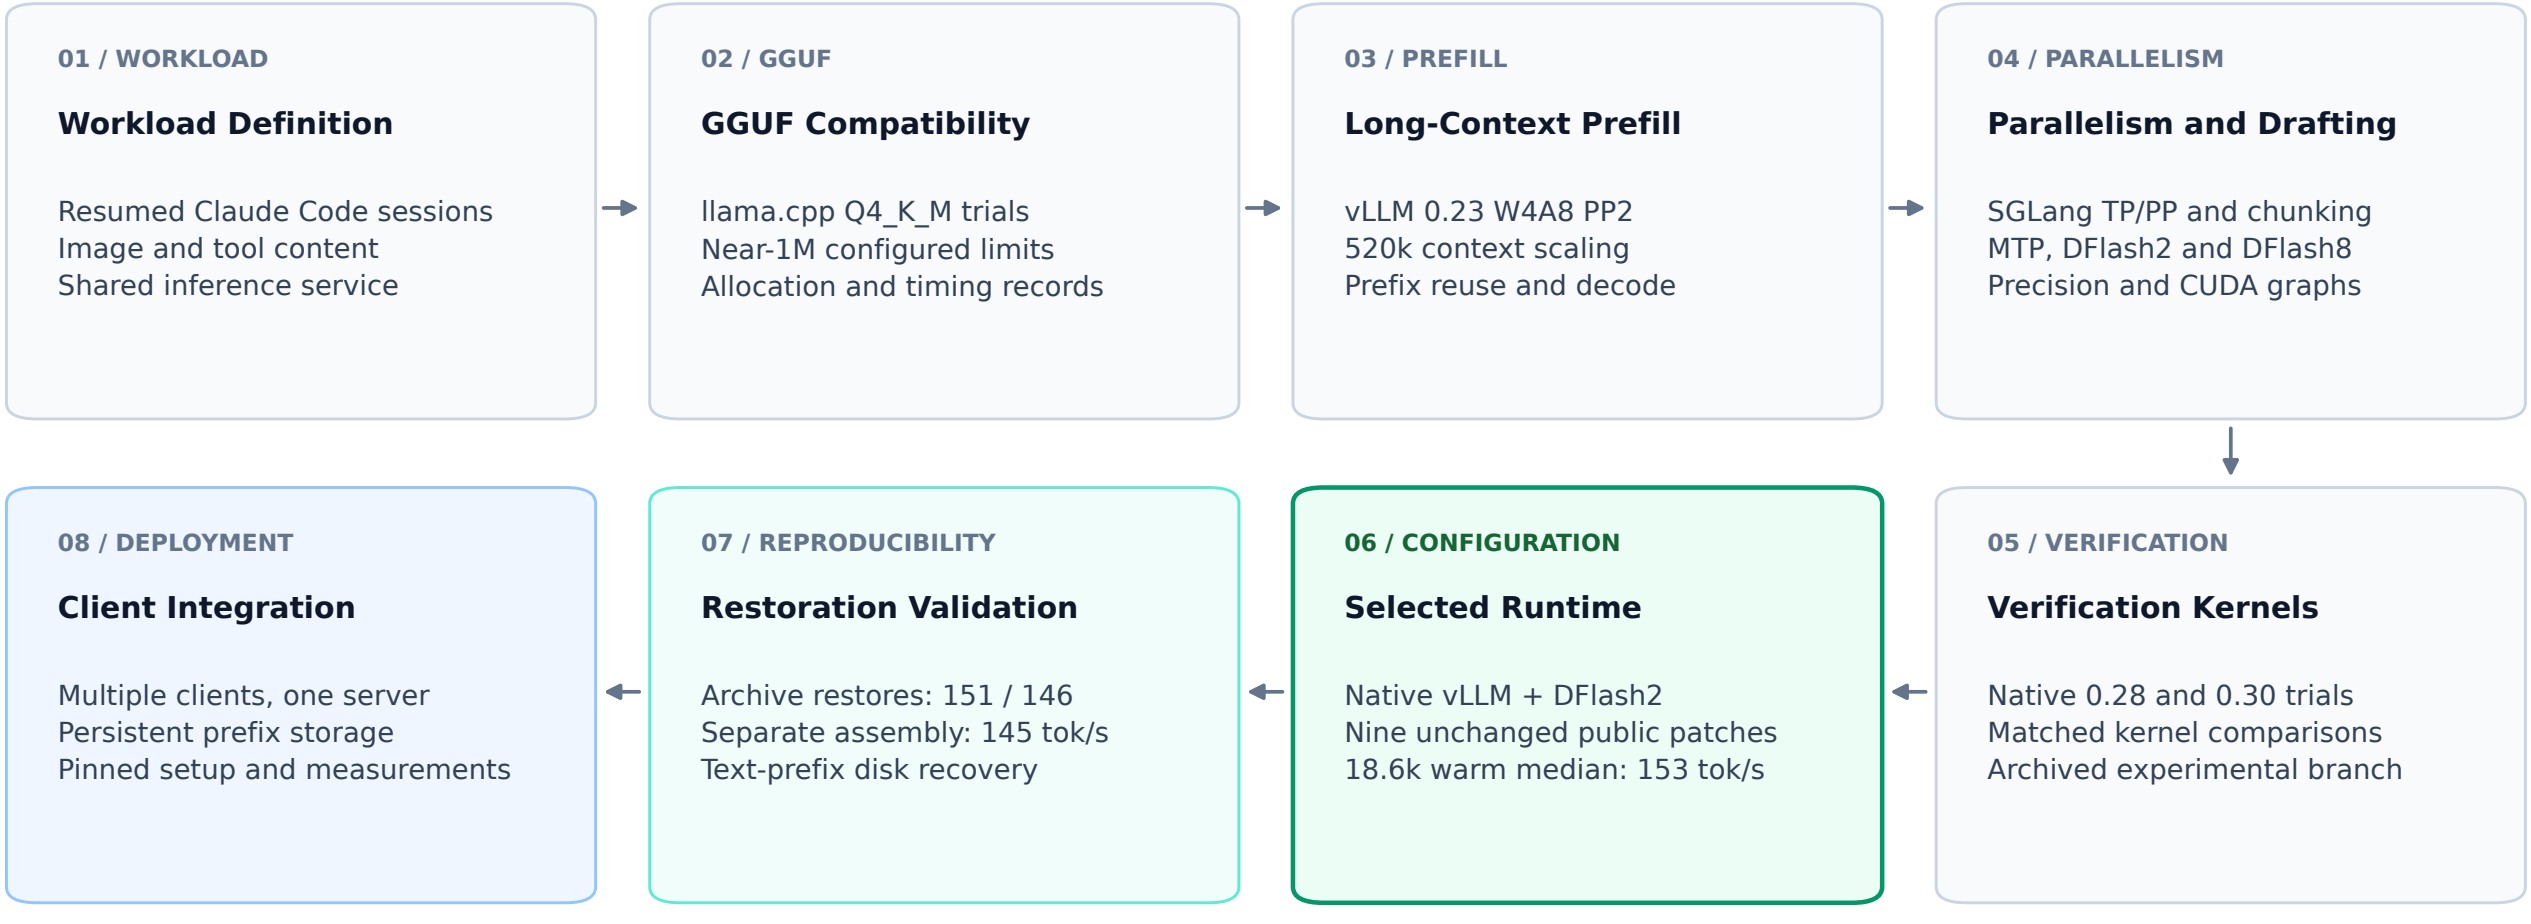

Campaigns overlap. Each result is linked to its configuration and original measurement record.

Measured headline: 18,649 actual input tokens with a warm prefix; output rates include reasoning.