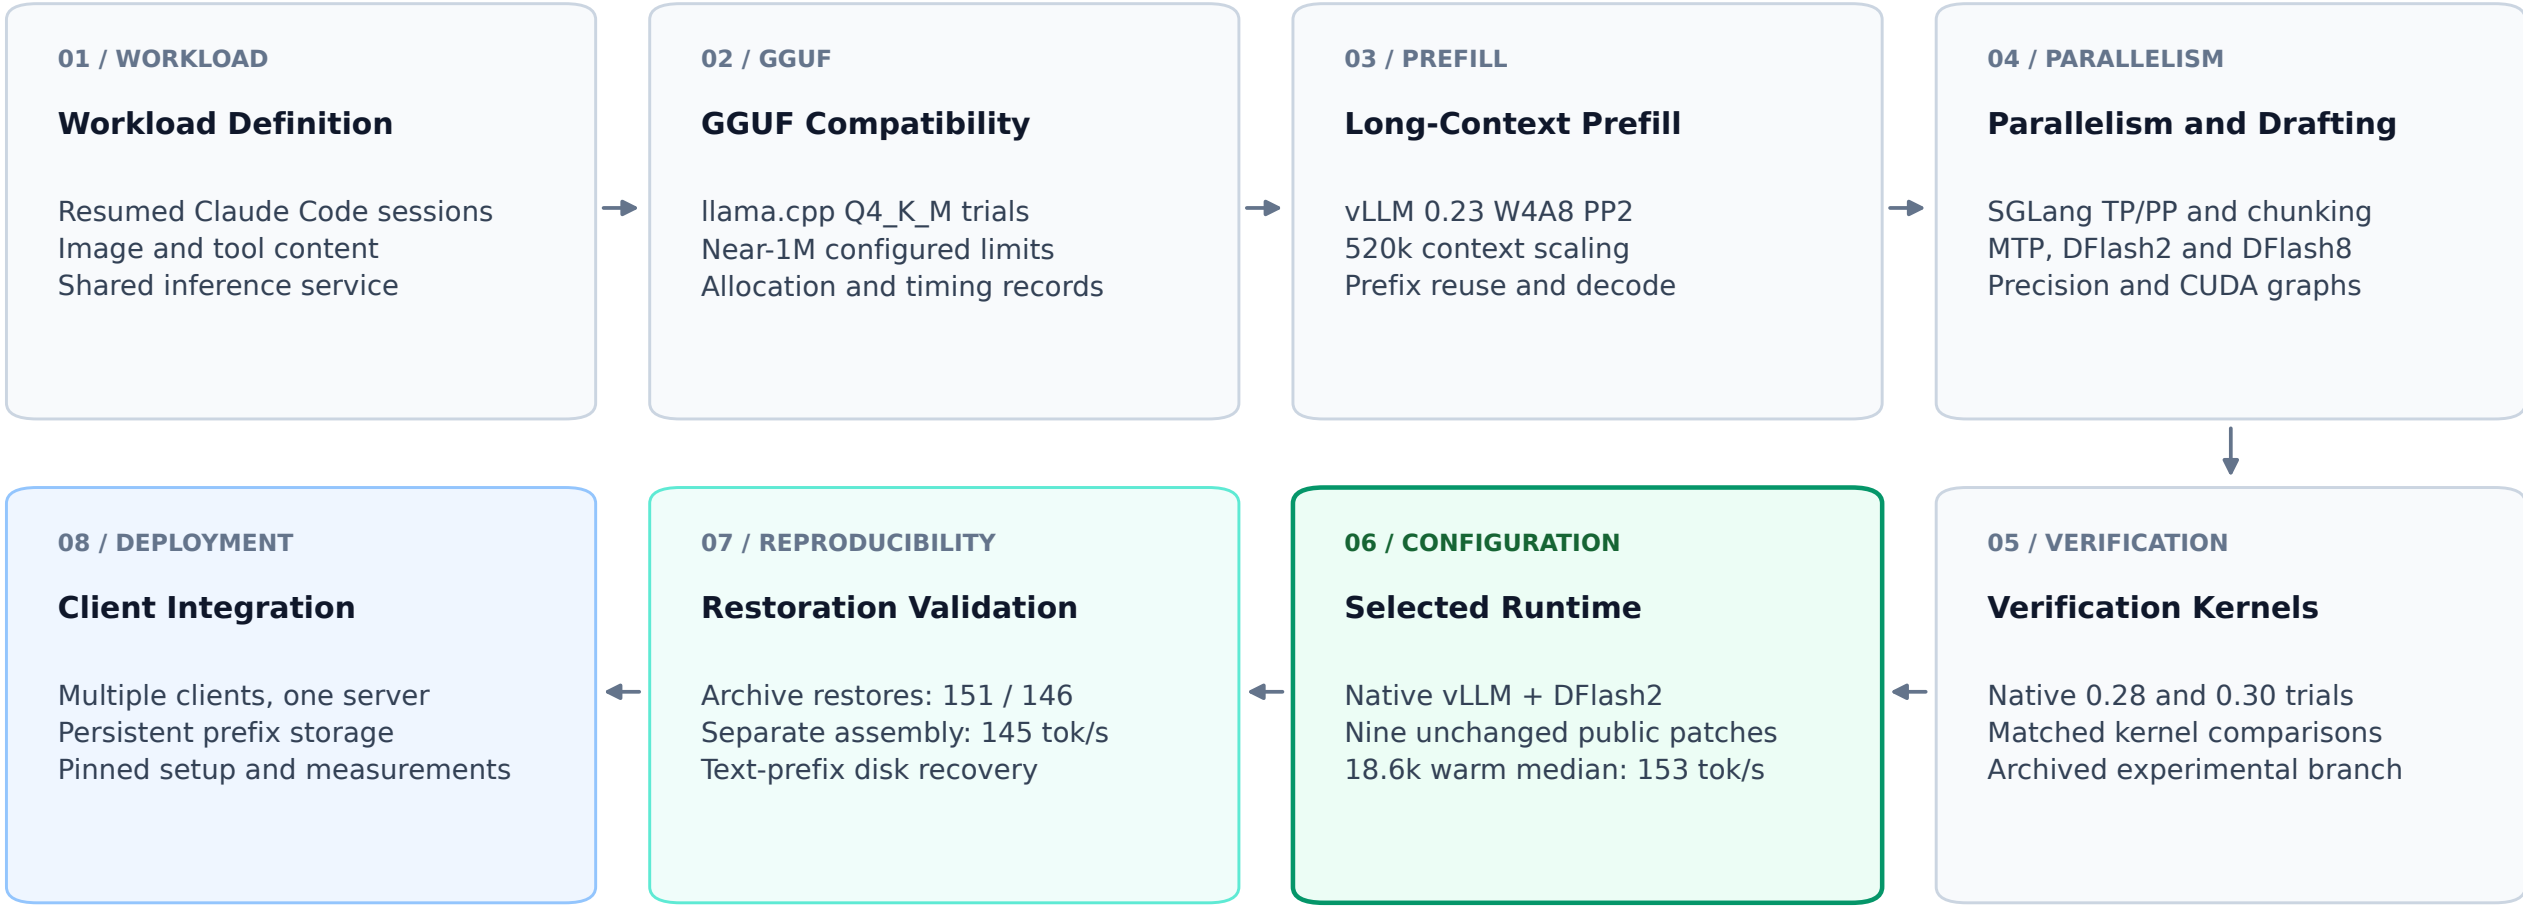
Dual-GPU Evaluation Overview

Engine compatibility, context scaling, speculative decoding and deployment validation

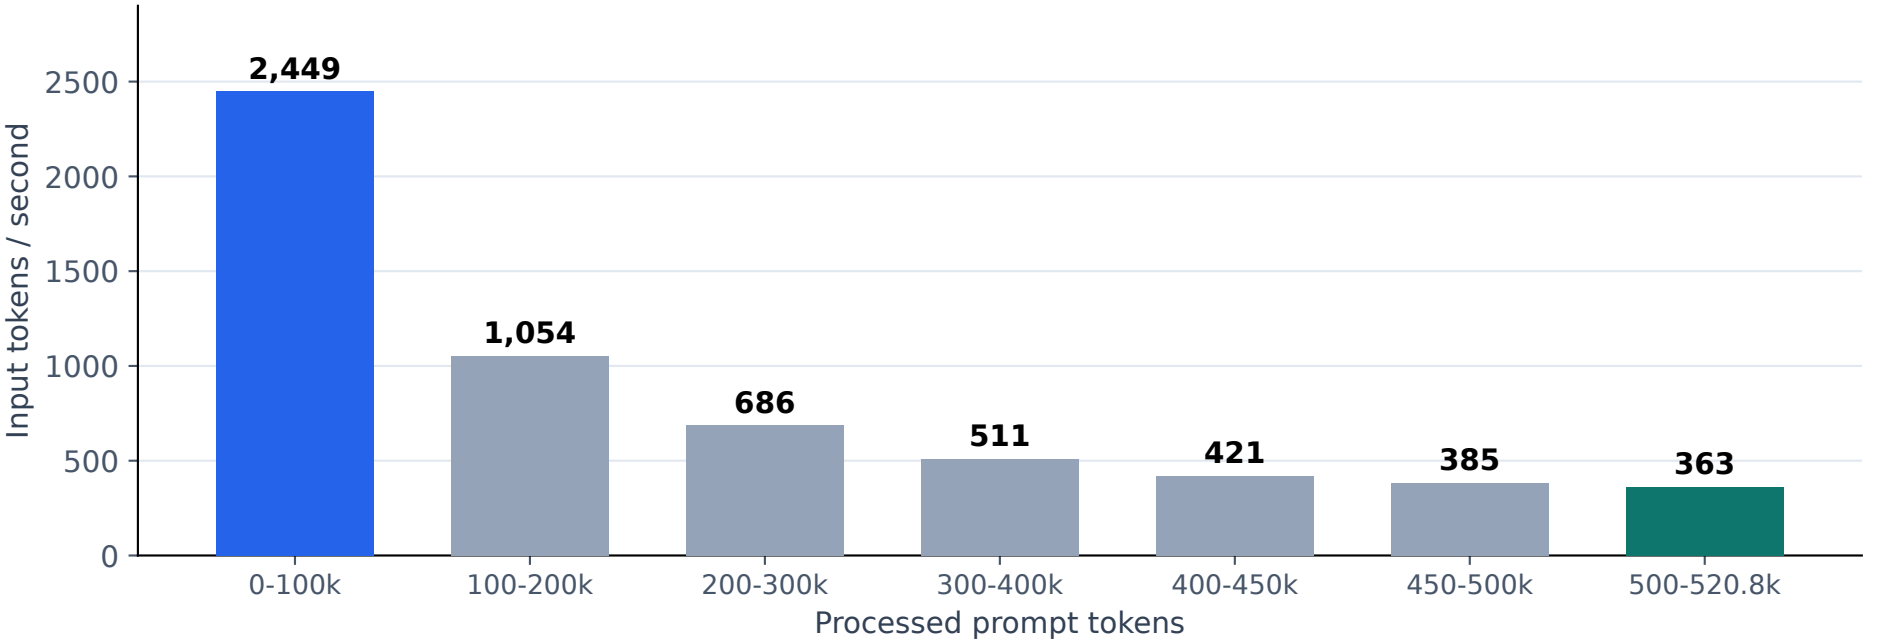


Campaigns overlap. Each result is linked to its configuration and original measurement record.

Measured headline: 18,649 actual input tokens with a warm prefix; output rates include reasoning.

Long-Context Prefill Scaling

vLLM 0.23 W4A8 PP2 | Median completed-batch prefill rate by processed-token range

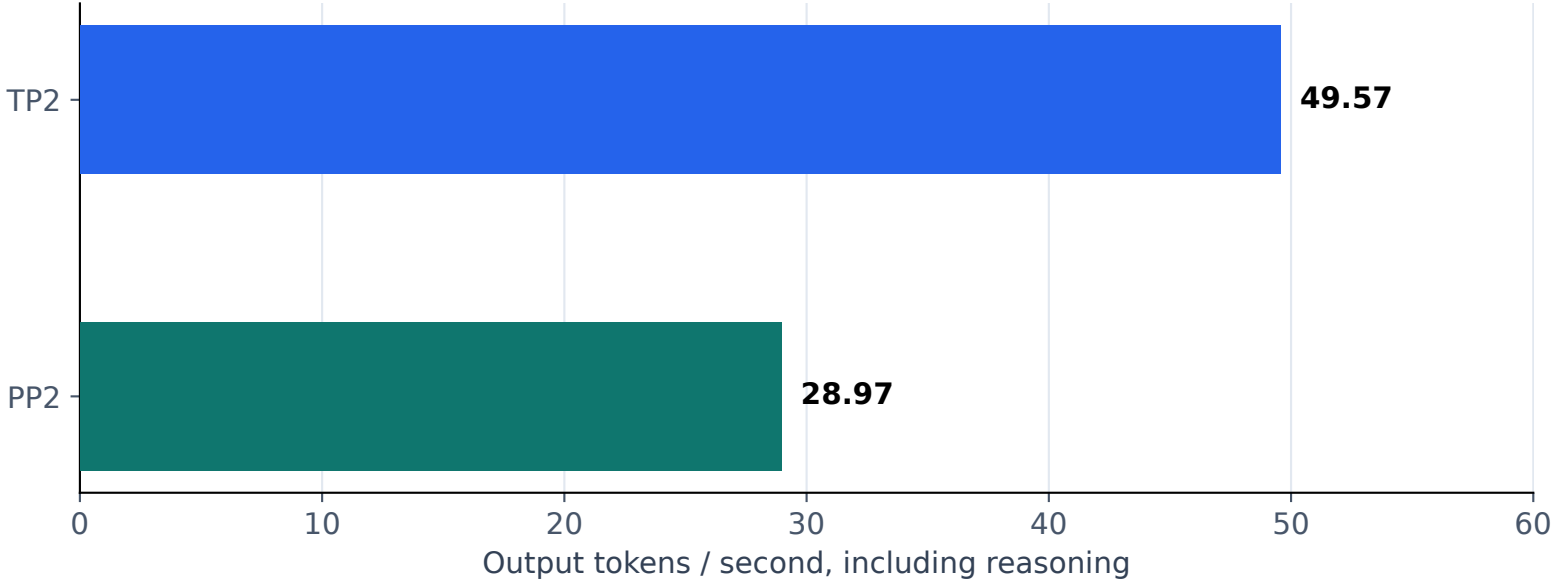


Actual input: 520,761 tokens. First batch excluded. Bins contain 31, 32, 32, 32, 16, 16 and 8 batches.

The separate 3,437 tok/s early sampling window uses a different metric. Source: S038 in benchmarks/history.

SGLang Parallelism Evaluation

SGLang 0.5.19 + MTP | 5,224 input / 4,096 output tokens | One run per layout

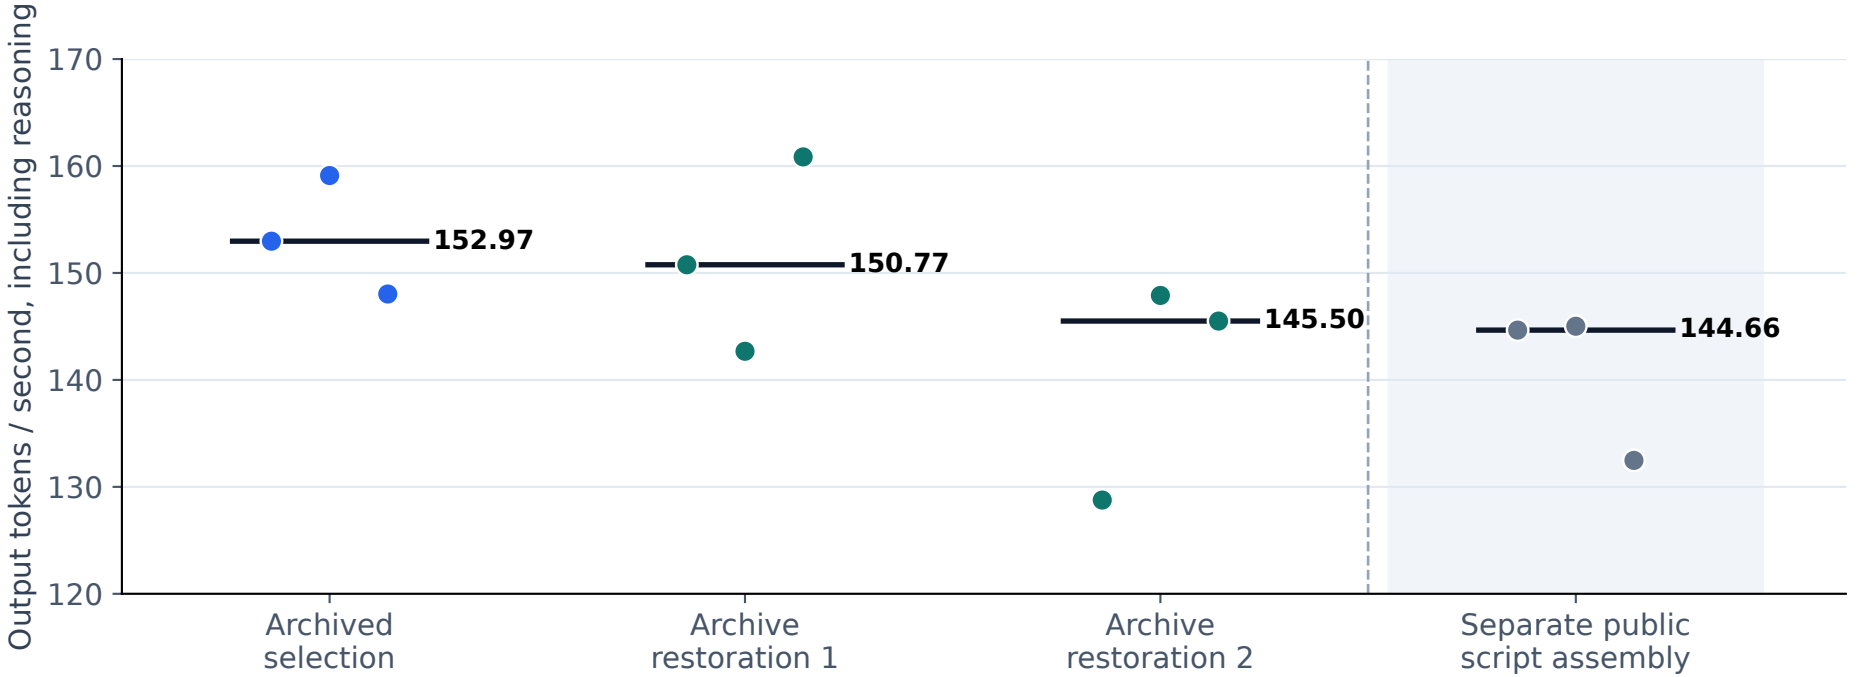


PP2: AWQ W4A16, FP8 KV, historical backports. Full TP2 launch recovery is incomplete.

First API phase only; later contexts diverge. Other configuration effects are not isolated. Source: S016.

Archive Restoration and Reconstruction

18,649 input / 1,024 output tokens per stream | Three observations in each group



Dots: individual capped runs, including reasoning. Black segments: group medians. Three observations per group.

First three groups: 17,920 prefix hits. Separate assembly: 16,128 hits, fixed 10.5 GiB KV/card, smaller shared pool.

Sources: benchmarks/results.json, history D008/D009/D010, and assembly-validation-20261017.json.