Long-Context Prefill Scaling

vLLM 0.23 W4A8 PP2 | Median completed-batch prefill rate by processed-token range

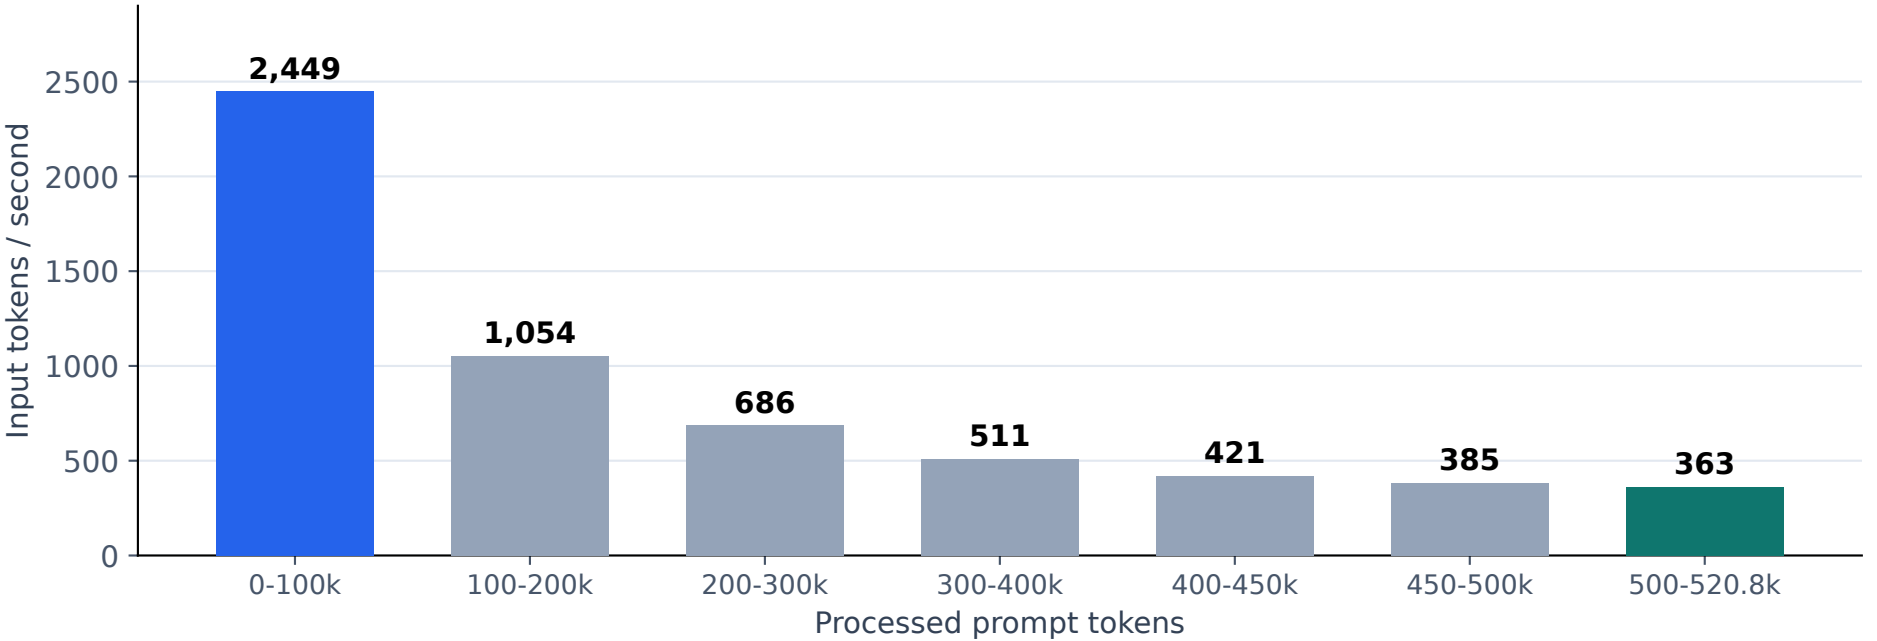

Actual input: 520,761 tokens. First batch excluded. Bins contain 31, 32, 32, 32, 16, 16 and 8 batches.

The separate 3,437 tok/s early sampling window uses a different metric. Source: S038 in benchmarks/history.